

CS8395 Assignment 0

Name: Daniel Yan

VUNet ID: yand1

Task 1

i).

```
Anaconda Prompt (2)

(base) C:\Users\Daniel>conda list anaconda
# packages in environment at F:\programs\anaconda:
#
# Name                        Version      Build      Channel
#-----
anaconda                      2019.03      py37_0
anaconda-client               1.7.2        py37_0
anaconda-navigator            1.9.7        py37_0
anaconda-project              0.8.2        py37_0

(base) C:\Users\Daniel>
```

ii).

```
Anaconda Prompt (2)

Total: 25.9 MB

The following NEW packages will be INSTALLED:

certifi      conda-forge/win-64::certifi-2019.11.28-py36_0
pip          conda-forge/win-64::pip-19.3.1-py36_0
python      conda-forge/win-64::python-3.6.7-he025d50_1006
setuptools   conda-forge/win-64::setuptools-44.0.0-py36_0
vc          pkgs/main/win-64::vc-14.1-h0510ff6_4
vs2015_runtime pkgs/main/win-64::vs2015_runtime-14.16.27012-hf0eaf9b_1
wheel        conda-forge/win-64::wheel-0.33.6-py36_0
wincertstore conda-forge/win-64::wincertstore-0.2-py36_1003

Proceed ([y]/n)? y

Downloading and Extracting Packages
pip-19.3.1      1.9 MB |#####| 100%
wheel-0.33.6    52 KB |#####| 100%
vs2015_runtime-14.15 2.4 MB |#####| 100%
python-3.6.7    20.7 MB |#####| 100%
setuptools-44.0.0 661 KB |#####| 100%
certifi-2019.11.28 148 KB |#####| 100%
wincertstore-0.2 13 KB |#####| 100%
Preparing transaction: done
Verifying transaction: done
Executing transaction: done
#
# To activate this environment, use
#
#     $ conda activate python36
#
# To deactivate an active environment, use
#
#     $ conda deactivate

(base) C:\Users\Daniel>conda activate python36

(python36) C:\Users\Daniel>
```



Task 2

```
Anaconda Prompt (2) - conda install pytorch torchvision cudatoolkit=10.1 -c pytorch
ninja          conda-forge/win-64::ninja-1.9.0-h1ad3211_1
numpy          conda-forge/win-64::numpy-1.17.3-py36hc71023c_0
olefile        conda-forge/noarch::olefile-0.46-py_0
pillow         conda-forge/win-64::pillow-7.0.0-py36h9ea1dd6_0
pyparser       conda-forge/win-64::pyparser-2.19-py36_1
pytorch        torch/win-64::pytorch-1.3.1-py3.6_cuda101_cudnn7_0
six            conda-forge/win-64::six-1.13.0-py36_0
tk             conda-forge/win-64::tk-8.6.10-hfa6e2cd_0
torchvision    pytorch/win-64::torchvision-0.4.2-py36_cu101
xz             conda-forge/win-64::xz-5.2.4-h2fa13f4_1001
zlib           conda-forge/win-64::zlib-1.2.11-h2fa13f4_1006
zstd           conda-forge/win-64::zstd-1.4.4-hd8ae53_1

Proceed ([y]/n)? y

Downloading and Extracting Packages
pytorch-1.3.1          | 480.3 MB | ##### | 100%
Preparing transaction: done
Verifying transaction: done
Executing transaction: done

(python36) C:\Users\Daniel>python -c
Argument expected for the -c option
usage: python [option] ... [-c cmd | -m mod | file | -] [arg] ...
Try 'python -h' for more information.

(python36) C:\Users\Daniel>"import torch; print(torch.__version__)"pyh

(python36) C:\Users\Daniel>python -c "import torch; print(torch.__version__)"
1.3.1

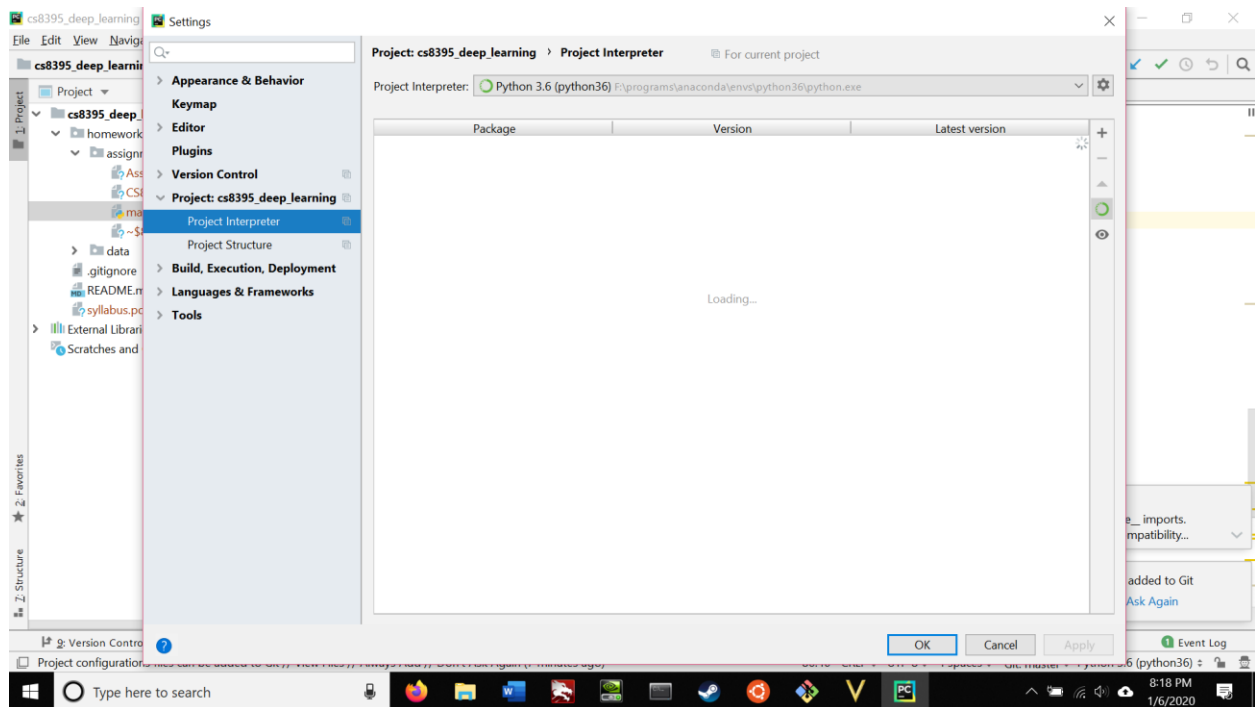
(python36) C:\Users\Daniel>
```

Task 3


i).

The image shows a PyCharm IDE window with a Python file named 'main.py' open. The file is located at 'F:\Daniel\Documents\vepos\cs8395_deep_learning\homework\assignment0\main.py'. The code defines a neural network class 'Net' that inherits from 'nn.Module'. The 'forward' method takes an input 'x' and processes it through several layers: a convolutional layer (conv1), a ReLU activation, another convolutional layer (conv2), a max pooling layer, a dropout layer, a fully connected layer (fc1), another ReLU activation, and a final fully connected layer (fc2). The 'init' method initializes these layers with specific parameters. The IDE interface includes a sidebar with a project tree showing the file structure, a top toolbar with various actions, and a bottom status bar with version control and terminal information. Two notifications are visible on the right side of the IDE, one about Python version compatibility and another about project configuration files.

ii).



Task 4



Task 5 (Description as comments in Code)

```
# Imports for Pytorch
from __future__ import print_function
import argparse
import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
from torchvision import datasets, transforms
from torch.optim.lr_scheduler import StepLR

# Define the neural network
class Net(nn.Module):
    # Define the dimensions for each layer.
    def __init__(self):
        super(Net, self).__init__()
        # First convolutional layer has 1 input channel, 32 output channels,
        # a 3x3 square kernel, and a stride of 1.
        self.conv1 = nn.Conv2d(1, 32, 3, 1)
        # Second convolutional layer has 32 input channels
        # since the first layer has 32 output channels.
        # The second layer has 64 output channels, uses
        # a 3x3 square kernel, and has a stride of 1.
        self.conv2 = nn.Conv2d(32, 64, 3, 1)
        # Dropout is performed twice in the network,
        # with the first time set to 0.25 and the
        # second time set to 0.5.
        self.dropout1 = nn.Dropout2d(0.25)
        self.dropout2 = nn.Dropout2d(0.5)
        # Two fully connected layers. The input shape to the
        # first fully connected layer is 64x12x12 = 9216. This is
        # because the MNIST image is 28x28, so the first
        # convolutional layer changes it to 26x26 since the kernel
        # is 3x3. The second convolutional layer changes it to 24x24.
        # We then have a maxpool layer that changes
        # the dimensions to 12x12. Since we have 64 channels as
        # the output from the second convolutional layer,
        # we get a total of 64x12x12 = 9216. The output from
        # the first fully connected layer is size 128.
        self.fc1 = nn.Linear(9216, 128)
        # Second fully connected layer takes in shape of
        # 128 from the output of the first fully connected layer
        # and then has 10 outputs because we have 10 classes for MNIST.
        self.fc2 = nn.Linear(128, 10)

# Define the structure for forward propagation.
def forward(self, x):
    # We begin with a convolutional layer with a
    # Relu activation function. We then use a second
    # convolutional layer and perform max pooling
    # and dropout on the output. We then flatten the
    # 64 channels from the output of the second
    # convolutional layer to pass to the first fully
    # connected layer, and use a Relu activation
    # function for the output. We then perform dropout
    # a second time and send the output for the
    # softmax function, since we are performing classification.
    x = self.conv1(x)
    x = F.relu(x)
    x = self.conv2(x)
    x = F.max_pool2d(x, 2)
    x = self.dropout1(x)
```

```

        x = torch.flatten(x, 1)
        x = self.fc1(x)
        x = F.relu(x)
        x = self.dropout2(x)
        x = self.fc2(x)
        output = F.log_softmax(x, dim=1)
        return output

def train(args, model, device, train_loader, optimizer, epoch):
    # Specify that we are in training phase
    model.train()
    # Iterate through all minibatches.
    for batch_idx, (data, target) in enumerate(train_loader):
        # Send training data and the training labels to GPU/CPU
        data, target = data.to(device), target.to(device)
        # Zero the gradients carried over from previous step
        optimizer.zero_grad()
        # Obtain the predictions from forward propagation
        output = model(data)
        # Compute the negative log likelihood of the loss function
        loss = F.nll_loss(output, target)
        # Perform backward propagation to compute the negative gradient, and
        # update the gradients with optimizer.step()
        loss.backward()
        optimizer.step()
        # Send output to log if logging is needed
        if batch_idx % args.log_interval == 0:
            print('Train Epoch: {} [{}/{}] ( {:.0f}%) \t Loss: {:.6f}'.format(
                epoch, batch_idx * len(data), len(train_loader.dataset),
                100. * batch_idx / len(train_loader), loss.item()))

def test(args, model, device, test_loader):
    # Specify that we are in evaluation phase
    model.eval()
    # Set the loss and number of correct instances initially to 0.
    test_loss = 0
    correct = 0
    # No gradient calculation because we are in testing phase.
    with torch.no_grad():
        # For each testing example, we run forward
        # propagation to calculate the
        # testing prediction. Update the total loss
        # and the number of correct predictions
        # with the counters from above.
        for data, target in test_loader:
            data, target = data.to(device), target.to(device)
            output = model(data)
            test_loss += F.nll_loss(output, target, reduction='sum').item()
            pred = output.argmax(dim=1, keepdim=True)
            correct += pred.eq(target.view_as(pred)).sum().item()

    # Average the loss by dividing by the total number of testing instances.
    test_loss /= len(test_loader.dataset)

    # Print out the statistics for the testing set.
    print('\nTest set: Average loss: {:.4f}, Accuracy: {}/{ } ( {:.0f}%) \n'.format(
        test_loss, correct, len(test_loader.dataset),
        100. * correct / len(test_loader.dataset)))

def main():

```

```

# Command line arguments for hyperparameters of
# training and testing batch size, the number of
# epochs, the learning rate, gamma, and other
# settings such as whether to use a GPU device, the
# random seed, how often to log, and
# whether we should save the model.
parser = argparse.ArgumentParser(description='PyTorch MNIST Example')
parser.add_argument('--batch-size', type=int, default=64, metavar='N',
                    help='input batch size for training (default: 64)')
parser.add_argument('--test-batch-size', type=int, default=1000, metavar='N',
                    help='input batch size for testing (default: 1000)')
parser.add_argument('--epochs', type=int, default=14, metavar='N',
                    help='number of epochs to train (default: 14)')
parser.add_argument('--lr', type=float, default=1.0, metavar='LR',
                    help='learning rate (default: 1.0)')
parser.add_argument('--gamma', type=float, default=0.7, metavar='M',
                    help='Learning rate step gamma (default: 0.7)')
parser.add_argument('--no-cuda', action='store_true', default=False,
                    help='disables CUDA training')
parser.add_argument('--seed', type=int, default=1, metavar='S',
                    help='random seed (default: 1)')
parser.add_argument('--log-interval', type=int, default=10, metavar='N',
                    help='how many batches to wait before logging training
status')

parser.add_argument('--save-model', action='store_true', default=False,
                    help='For Saving the current Model')
args = parser.parse_args()
# Command to use gpu depending on command line arguments and if there is a cuda
device
use_cuda = not args.no_cuda and torch.cuda.is_available()

# Random seed to use
torch.manual_seed(args.seed)

# Set to either use gpu or cpu
device = torch.device("cuda" if use_cuda else "cpu")

# GPU keywords.
kwargs = {'num_workers': 1, 'pin_memory': True} if use_cuda else {}
# Load in the training and testing datasets. Convert to
# pytorch tensor and normalize.
train_loader = torch.utils.data.DataLoader(
    datasets.MNIST('./data', train=True, download=True,
                  transform=transforms.Compose([
                      transforms.ToTensor(),
                      transforms.Normalize((0.1307,), (0.3081,))
                  ])),
    batch_size=args.batch_size, shuffle=True, **kwargs)
test_loader = torch.utils.data.DataLoader(
    datasets.MNIST('./data', train=False, transform=transforms.Compose([
        transforms.ToTensor(),
        transforms.Normalize((0.1307,), (0.3081,))
    ])),
    batch_size=args.test_batch_size, shuffle=True, **kwargs)

# Run model on GPU if available
model = Net().to(device)
# Specify Adadelta optimizer
optimizer = optim.Adadelta(model.parameters(), lr=args.lr)

# Run for the set number of epochs. For each epoch, run the training
# and the testing steps. Scheduler is used to specify the learning rate.

```

```
scheduler = StepLR(optimizer, step_size=1, gamma=args.gamma)
for epoch in range(1, args.epochs + 1):
    train(args, model, device, train_loader, optimizer, epoch)
    test(args, model, device, test_loader)
    scheduler.step()

# Save model if specified by the command line argument
if args.save_model:
    torch.save(model.state_dict(), "mnist_cnn.pt")

if __name__ == '__main__':
    main()
```